

# Rules of Thumb for Activation Functions in Neural Networks and Matrix Shape Notes

## Rules of Thumb for Activation Functions in Neural Networks

### Overview

Activation functions are essential in neural networks. They introduce non-linearity, help models learn complex patterns, and influence training stability. Below are the key rules of thumb.

### General Guidelines

- Avoid Sigmoid in deep networks – can cause vanishing gradients and slow convergence; still useful in the output layer for binary classification.
- ReLU is usually the default choice – simple, efficient, reduces vanishing gradient, but may suffer from “dying ReLU” problem.
- Leaky ReLU or variants (PReLU, ELU) help when ReLU neurons die, since they allow small gradients for negative inputs.
- Use Tanh if negative outputs are useful – zero-centered but still prone to vanishing gradients in deeper nets.
- Softmax is for multi-class classification outputs – converts logits into probability distribution across classes.
- Linear activation for regression tasks – no activation in the final layer if predicting continuous values.

### Practical Rules of Thumb

- Hidden layers: Start with ReLU; try Leaky ReLU/ELU if training is unstable.
- Binary classification output: Use sigmoid in the final layer.

- Multi-class classification output: Use softmax in the final layer.
- Regression output: Use linear activation (no activation) in the final layer.
- Small/shallow networks: Tanh may sometimes work better.
- RNNs: Typically use tanh or sigmoid inside cells; ReLU in newer architectures.
- Normalization: Combine batch normalization with ReLU for more stable training.

### Quick Mnemonic

- ReLU = default for hidden layers
- Sigmoid = binary output
- Softmax = multi-class output
- Linear = regression output
- Tanh = centered data / small networks

## General Guidelines (Detailed)

### Activation Function Guidelines

- **Sigmoid:** Range:  $(0, 1)$ ; used in binary output. Decision rule: if output  $\geq 0.5 \rightarrow$  class 1, else class 0. Problems: vanishing gradients, outputs not zero-centered.
- **Tanh:** Range:  $(-1, 1)$ ; zero-centered and generally better than sigmoid. Useful in shallow networks or RNNs but still prone to vanishing gradients in deep networks.
- **ReLU (Rectified Linear Unit):** Range:  $[0, \infty)$ ; default activation for hidden layers. Fast and efficient; helps reduce vanishing gradients. Risk: “dying ReLU” problem — neurons stuck at zero output.
- **Leaky ReLU / PReLU / ELU:** Range:  $(-\infty, \infty)$ ; allows a small non-zero gradient for negative inputs. Solves the dead neuron issue in standard ReLU. Recommended when ReLU performance stagnates.
- **Softmax:** Range:  $(0, 1)$  for each output neuron, with all outputs summing to 1. Converts logits into a probability distribution. Used exclusively in the output layer for multi-class classification.
- **Linear (Identity Function):** Range:  $(-\infty, \infty)$ ; used in regression outputs. No squashing effect — allows real-valued predictions without restriction.

## Activation Function Comparison Table

Activation	Value Range	Typical Case	Use	Notes
Sigmoid ( $\sigma$ )	(0, 1)	Binary classification output		Threshold at 0.5 → label 1 else 0; not good for hidden layers (vanishing gradients).
Tanh	(−1, 1)	Shallow networks, RNNs		Zero-centered; better than sigmoid but still vanishes in deep networks.
ReLU	[0, $\infty$ )	Default for hidden layers		Fast, efficient, but risk of “dying ReLU.”
Leaky ReLU / PReLU / ELU	(− $\infty$ , $\infty$ )	Hidden layers when ReLU neurons die		Allows small gradient for negative inputs; mitigates dead neurons.
Softmax	(0, 1), sum = 1	Multi-class classification output		Converts logits into a normalized probability distribution.
Linear	(− $\infty$ , $\infty$ )	Regression outputs		No squashing; direct prediction of real values.

## Matrix Shape, Dimensions, and Neural Network Equation

### 1. Shape of a Matrix

In NumPy (and mathematics generally): `.shape` → gives the size of each dimension as a tuple.

Listing 1: Matrix shape example

```
A = [[1,2,3],
     [4,5,6]]
A.shape # (2,3)    2 rows   3 columns
```

`shape[0]` → number of rows. `shape[1]` → number of columns.

## 2. Dimensions (Rank of Array)

- Scalar: no dimension  $\rightarrow ()$
- Vector: 1D  $\rightarrow (n,)$
- Matrix: 2D  $\rightarrow (m,n)$
- Tensor: higher-D  $\rightarrow$  e.g., (batch, height, width, channels)

## 3. Matrix Operations

### Matrix Operations Summary

**Addition/Subtraction:** Two matrices can be added/subtracted only if shapes match. Example:  $(m,n) + (m,n) \rightarrow (m,n)$

### Multiplication:

1. Elementwise (Hadamard):  $(m,n)*(m,n) \rightarrow (m,n)$
2. Matrix multiplication: If A is  $(m,n)$  and B is  $(n,p)$ , then  $C=A \cdot B \rightarrow (m,p)$ .  
Rule: Inner dimensions must match.

## 4. Neural Network Equation

Forward pass:  $z = W \cdot x + b$

$x \rightarrow (n, 1)$	input vector (n features)
$W \rightarrow (m, n)$	weight matrix (m neurons, n inputs)
$b \rightarrow (m, 1)$	bias vector
$z \rightarrow (m, 1)$	output

$$W(m, n) \times x(n, 1) \rightarrow (m, 1), \quad +b(m, 1) \rightarrow (m, 1)$$

### Example

Input has 4 features  $\rightarrow x.shape = (4, 1)$  Layer has 3 neurons  $\rightarrow z.shape = (3, 1)$  Then:

$$W.shape = (3, 4), \quad b.shape = (3, 1), \quad z = W \cdot x + b \rightarrow (3, 1)$$

Symbol	Shape	Meaning
x	(n,1)	Input features
W	(m,n)	Weights
b	(m,1)	Bias
z	(m,1)	Linear output

# NumPy .shape Attribute – Notes with Examples

## 1. Overview

Every ndarray object in NumPy has a `.shape` attribute that returns a tuple showing size along each axis.

## 2. 1D Arrays (Vectors)

```
import numpy as np
a = np.array([10, 20, 30, 40, 50])
print(a.shape)      # (5,)
print(a.ndim)      # 1
print(a.shape[0])  # 5
```

**Explanation:** 1D array with 5 elements → `.shape = (5,)`

## 3. 2D Arrays (Matrices)

```
b = np.array([[1, 2, 3],
              [4, 5, 6],
              [7, 8, 9],
              [10, 11, 12]])
print(b.shape) # (4, 3)
```

Explanation: 4 rows, 3 columns → `(4,3)`

## 4. 3D Arrays (Tensors)

```
c = np.array([
    [[1, 2], [3, 4], [5, 6]],
    [[7, 8], [9, 10], [11, 12]]
])
```

`c.shape = (2, 3, 2), c.ndim = 3`

## 5. 4D Arrays (e.g., Image Batches)

```
d = np.random.rand(10, 64, 64, 3)
print(d.shape) # (10, 64, 64, 3)
```

Explanation: 10 images,  $64 \times 64$ , 3 RGB channels.

## 6. Practical Uses

```
# Iteration by rows
rows, cols = b.shape
for i in range(rows):
    print("Row:", b[i])

# Reshaping arrays
arr = np.arange(12)
reshaped = arr.reshape(3, 4)
print(reshaped.shape) # (3, 4)
```

## 7. Quick Reference Table

Array Type	Example Shape	Meaning
1D vector	(5,)	5 elements
2D matrix	(4, 3)	4 rows, 3 columns
3D tensor	(2, 3, 2)	2 blocks, each $3 \times 2$
4D tensor	(10, 64, 64, 3)	10 images, $64 \times 64$ pixels, 3 channels